

# Report

Cuiyi

November 8, 2014

# Discriminate Method Review

## 1 Logistic Regression

Assume  $y = 0$  or  $y = 1$ , and  $h_w(x) = \frac{1}{1+e^{-w^T x}}$

## 2 Perceptron Learning Algorithm

$h_w(x)$  is a threshold function:

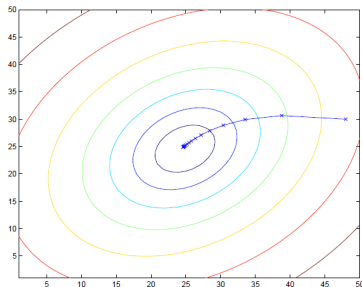
$$h_w(x) = g_w(w^T x) = \begin{cases} 1 & \text{if } w^T x \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

# Gradient descent Review

Batch Gradient descent

Repeat for each parameter :

$$w_i := w_i + \alpha \frac{\partial}{\partial w_i} \text{Loss}(h_w(x), y)$$



## Gradient descent Review(2)

The normal equations:

For a function  $f : \mathbb{R}^{m \times n} \rightarrow \mathbb{R}$  define the derivative of  $f$  with respect to  $A$  to be:

$$\nabla_A f(A) = \begin{pmatrix} \frac{\partial}{\partial A_{11}} & \cdots & \frac{\partial}{\partial A_{1n}} \\ \vdots & \ddots & \vdots \\ \frac{\partial}{\partial A_{m1}} & \cdots & \frac{\partial}{\partial A_{mn}} \end{pmatrix}$$

For examples, suppose  $A = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix}$  is a 2-by-2 matrix, and:

$$f(A) = \frac{3}{2}A_{11} + 5A_{12}^2 + A_{21}A_{22}$$

Then we have:

$$\nabla_A f(A) = \begin{pmatrix} \frac{3}{2} & 10A_{12} \\ A_{22} & A_{21} \end{pmatrix}$$

Define of **trace** operator. For an n-by-n matrix A:

$$tr A = \sum_{i=1}^n A_{ii}$$

## Gradient descent Review(3)

For two matrices  $A$  and  $B$ , such that  $ABC$  is square, we have:

$$\text{tr}AB = \text{tr}BA$$

As corollaries:

$$\text{tr}ABC = \text{tr}CAB = \text{tr}BCA$$

Here  $A$  and  $B$  is square matrices and  $a$  is real number:

$$\text{tr}A = \text{tr}A^T$$

$$\text{tr}(A + B) = \text{tr}A + \text{tr}B$$

$$\text{tr}aA = a\text{tr}A$$

## Gradient descent Review(4)

Some facts without proof:

$$\nabla_A \text{tr} AB = B^T$$

$$\nabla_{A^T} f(A) = (\nabla_A f(A))^T$$

$$\nabla_A \text{tr} ABA^T C = CAB + C^T AB^T$$

Revisit Least Squares:

define **design matrix**  $X$  to be the m-by-n matrix:

$$X = \begin{pmatrix} x_1^{(1)} & \dots & x_n^{(1)} \\ \vdots & \ddots & \vdots \\ x_1^{(m)} & \dots & x_n^{(m)} \end{pmatrix}$$

## Gradient descent Review(5)

define  $\vec{y}$  be m-dimensional vector:

$$\vec{y} = \begin{pmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{pmatrix}$$

Since  $h_w(x^{(i)}) = (x^{(i)})^T w$ :

$$\begin{aligned} Xw - \vec{y} &= \begin{pmatrix} (x^{(1)})^T w \\ \vdots \\ (x^{(m)})^T w \end{pmatrix} - \begin{pmatrix} y^{(1)} \\ \vdots \\ y^{(m)} \end{pmatrix} \\ &= \begin{pmatrix} h_w(x^{(1)}) - y^{(1)} \\ \vdots \\ h_w(x^{(m)}) - y^{(m)} \end{pmatrix} \end{aligned}$$

## Gradient descent Review(5)

Using the fact that for  $\vec{z}$ , we have  $\vec{z}^T \vec{z} = \sum_i z_i^2$ :

$$\begin{aligned}\frac{1}{2}(Xw - \vec{y})^T(Xw - \vec{y}) &= \frac{1}{2} \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)})^2 \\ &= J(w)\end{aligned}$$

$$\begin{aligned}\nabla_w J(w) &= \nabla_w \frac{1}{2} (Xw - \vec{y})^T (Xw - \vec{y}) \\ &= \frac{1}{2} \nabla_w \text{tr}(w^T X^T X w - w^T X^T \vec{y} - \vec{y}^T X w + \vec{y}^T \vec{y}) \\ &= \frac{1}{2} \nabla_w (\text{tr} w^T X^T X w - 2 \text{tr} \vec{y}^T X w) \\ &= \frac{1}{2} (X^T X w + X^T X w - 2 X^T \vec{y}) \\ &= X^T X w - X^T \vec{y}\end{aligned}$$



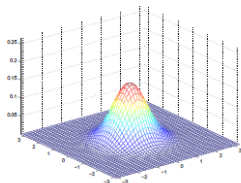
# What are Generative Learning algorithms

- Consider a classification problem: distinguish between benign tumors and malignant tumors.
- Given a training set, an algorithm like logistic regression or the perceptron algorithm (basically) tries to find a straight line that is, a decision boundary that separates the benign tumor and malignant tumor. Algorithms that try to learn  $P(y|x)$  directly are called **discriminative learning algorithms**.
- Different approach: First we can build a model of what benign tumors look like. Then we build a separate model of what malignant tumors look like. Finally, to classify a new sample, we can match the new sample with each model, to see whether one model matches better than one other model. Algorithms that try to learn  $P(x|y)$  (and  $P(y)$ ) are called **generative learning algorithms**.

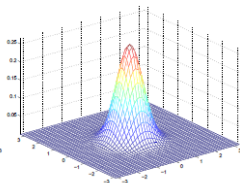
# The multivariate normal distribution

Suppose there is a multivariate normal distribution is parameterized by a **mean vector**  $\mu \in \mathbb{R}^n$  and a **covariance matrix**  $\Sigma \in \mathbb{R}^{n \times n}$ , its density is given by:

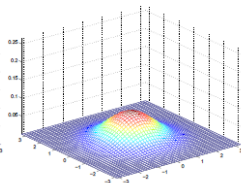
$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left( -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$$



$$\Sigma = I$$

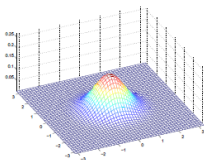


$$\Sigma = 0.6I;$$

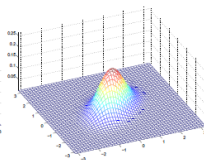


$$\Sigma = 2I.$$

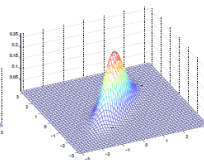
# The multivariate normal distribution(2)



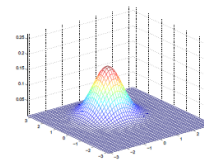
$$\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$



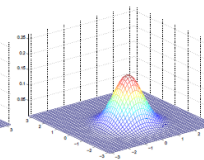
$$\Sigma = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix}$$



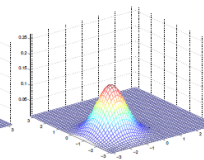
$$\Sigma = \begin{bmatrix} 1 & 0.8 \\ 0.8 & 1 \end{bmatrix}$$



$$\mu = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$



$$\mu = \begin{bmatrix} -0.5 \\ 0 \end{bmatrix}$$



$$\mu = \begin{bmatrix} -1 \\ -1.5 \end{bmatrix}$$

# The Gaussian Discriminant Analysis model(1)

GDA Model is:

$$\begin{aligned}y &\sim \text{Bernoulli}(\phi) \\x|y=0 &\sim \mathcal{N}(\mu_0, \Sigma) \\x|y=1 &\sim \mathcal{N}(\mu_1, \Sigma)\end{aligned}$$

Writing out the distribution:

$$p(y) = \phi^y(1 - \phi)^{1-y}$$

$$p(x|y=0) = \frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_0)^T \Sigma^{-1}(x - \mu_0)\right)$$

$$p(x|y=1) = \frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_1)^T \Sigma^{-1}(x - \mu_1)\right)$$

# The Gaussian Discriminant Analysis model(2)

The log-likelihood is:

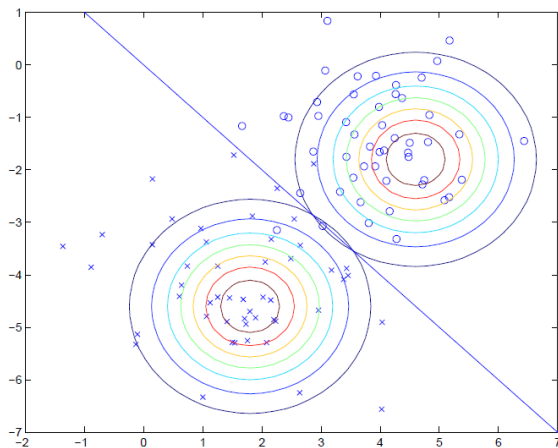
$$\begin{aligned}l(\phi, \mu_0, \mu_1, \Sigma) &= \log \prod_{i=1}^m p(x^{(i)}, y^{(i)}; \phi, \mu_0, \mu_1, \Sigma) \\&= \log \prod_{i=1}^m p(x^{(i)} | y^{(i)}; \mu_0, \mu_1, \Sigma) p(y^{(i)}; \phi)\end{aligned}$$

By maximizing  $l$  with respect to the parameters, we find the maximum likelihood estimate of parameters to be:

$$\begin{aligned}\phi &= \frac{1}{m} \sum_{i: y^{(i)}=1} y^{(i)} \\ \mu_0 &= \frac{\sum_{i: y^{(i)}=0} x^{(i)}}{\sum_{i: y^{(i)}=0} y^{(i)}} \\ \mu_1 &= \frac{\sum_{i: y^{(i)}=1} x^{(i)}}{\sum_{i: y^{(i)}=1} y^{(i)}} \\ \Sigma &= \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{y^{(i)}})(x^{(i)} - \mu_{y^{(i)}})^T\end{aligned}$$

## The Gaussian Discriminant Analysis model(3)

Pictorially, what the algorithm is doing can be seen in as follows:



## Discussion: GDA and logistic regression

If we view  $p(y = 1|x; \phi, \mu_0, \mu_1, \Sigma)$  as a function of  $x$ , there is a fact that:

$$p(y = 1|x; \phi, \Sigma, \mu_0, \mu_1) = \frac{1}{1 + \exp(-w^T x)}$$

This is exactly the form that logistic regression used to model  $p(y = 1|x)$ .

In fact in binary division problem:

$$p(x|y = 1) \sim \text{ExpFamily}(x, \eta) \Rightarrow p(y = 1|x) = \text{Sigmoid}$$

When would we prefer one model over another?

- GDA makes stronger modelling assumptions, and is more data efficient (i.e., requires less training data to learn well) when the modeling assumptions are correct or at least approximately correct.
- Logistic regression makes weaker assumptions, and is significantly more robust to deviations from modeling assumptions.

## Bayes' formula

Suppose that we have the prior probability  $P(\omega_j)$  and the conditional densities  $p(x|\omega_j)$ , the posterior probability can be written as:

$$P(\omega_j|x) = \frac{p(x|\omega_j)P(\omega_j)}{p(x)},$$

where

$$p(x) = \sum_{j=1}^c p(x|\omega_j)P(\omega_j).$$

Bayes' formula can be expressed informally in English:

$$\textit{posterior} = \frac{\textit{likelihood} \times \textit{prior}}{\textit{evidence}}.$$



## Bayesian Decision Theory Continuous Features

Let  $\omega_1, \dots, \omega_c$  be the finite set of  $c$  states of nature ("categories") and  $\alpha_1, \dots, \alpha_a$  be the finite set of a possible action. The Loss function  $\lambda(\alpha_i|\omega_j)$  describe the loss incurred for taking action  $\alpha_i$  when the state of nature is  $\omega_j$ . The posterior probability  $P(\omega_j|x)$  can be compute from  $P(\omega_j|x)$  by Bayes'formula:

$$P(\omega_j|x) = \frac{p(x|\omega_j)P(\omega_j)}{p(x)}$$

where the evidence is now

$$p(x) = \sum_{j=1}^c p(x|\omega_j)P(\omega_j)$$

## Bayesian Decision Theory Continuous Features(2)

Suppose that we observe a particular  $x$  and that we contemplate taking action  $\alpha_i$ .

$$R(\alpha_i|x) = \sum_{j=1}^c \lambda(\alpha_i|\omega_j)P(\omega_j|x)$$

In decision-theoretic terminology, an expected loss is called a risk, and  $R(\alpha_i|x)$  is called the **conditional risk**. If we define loss function to be

$$\lambda(\alpha_i|\omega_j) = \begin{cases} 0 & i = j \\ 1 & i \neq j \end{cases} \quad i, j = 1, \dots, c.$$

Condition risk will be

$$R(\alpha_i|x) = \sum_{j=1}^c \lambda(\alpha_i|\omega_j)P(\omega_j|x) = 1 - P(\omega_i|x)$$

# Classifiers, Discriminant Functions and Decision Surfaces

For 0-1 Loss function case, we can simply let the discriminant functions to be  $g_i(x) = P(\omega_i|x)$ . More generally, if we replace every  $g_i(x)$  by  $f(g_i(x))$ , where  $f(\cdot)$  is a monotonically increasing function, the resulting classification is unchanged.

$$g_i(x) = P(\omega_i|x) = \frac{p(x|\omega_i)P(\omega_i)}{\sum_{j=1}^c p(x|\omega_j)P(\omega_j)}$$

$$g_i(x) = p(x|\omega_i)P(\omega_i)$$

$$g_i(x) = \log p(x|\omega_i) + \log P(\omega_i)$$

# Discriminant Functions for the Normal Density

if  $p(x|\omega_i) \sim \mathcal{N}(\mu_i, \Sigma_i)$ ,

$$g_i(x) = -\frac{1}{2}(x - \mu_i)^T \Sigma_i^{-1} (x - \mu_i) - \frac{d}{2} \log 2\pi - \frac{1}{2} \log |\Sigma_i| + \log P(\omega_i)$$

Case:  $\Sigma_i = \sigma^2 I$

Ignore terms which are independent of  $i$ , we obtain:

$$\begin{aligned} g_i(x) &= -\frac{\|x - \mu_i\|^2}{2\sigma^2} + \log P(\omega_i) \\ &= -\frac{1}{2\sigma^2} (x^T x - 2\mu_i^T x + \mu_i^T \mu_i) + \log P(\omega_i) \end{aligned}$$

Since  $x^T x$  is the same for all  $i$ 's, we obtain

$$g_i(x) = w_i^T x + w_{i0}$$

where

$$w_i = \frac{1}{\sigma^2} \mu_i, \quad w_{i0} = -\frac{1}{2\sigma^2} \mu_i^T \mu_i + \log P(\omega_i)$$

## Discriminant Functions for the Normal Density(2)

For two class, we can set  $g_i(x) = g_j(x)$ :

$$w^T(x - x_0) = 0,$$

where,

$$w = \mu_i - \mu_j$$

and,

$$x_0 = \frac{1}{2}(\mu_i + \mu_j) - \frac{\sigma^2}{\|\mu_i - \mu_j\|} \log \frac{P(\omega_i)}{P(\omega_j)} (\mu_i - \mu_j)$$

## Discriminant Functions for the Normal Density(3)

Case  $\Sigma_i = \Sigma$ :

$$g_i(x) = w^T x + w_0$$

where

$$w_i = \Sigma^{-1} \mu_i, \quad w_{i0} = -\frac{1}{2} \mu_i^T \Sigma^{-1} \mu_i + \log P(\omega_i)$$

Case  $\Sigma_i = ?$

$$g_i = x^T W_i x + w_i^T x + w_{i0}$$

where,

$$W_i = -\frac{1}{2} \Sigma_i^{-1}, \quad w_i = \Sigma_i^{-1} \mu_i$$

and,

$$w_{i0} = -\frac{1}{2} \mu_i^T \Sigma_i^{-1} \mu_i - \frac{1}{2} \log |\Sigma_i| + \log P(\omega_i)$$

## Bayesian estimation

In Bayesian estimation, we consider parameter  $\theta$  to be a random variable. Any information we might have about  $\theta$  prior to observing the samples is assumed to be contained in a *known* prior density  $p(\theta)$ . Observation of the samples converts this to a posterior density  $p(\theta|\mathcal{D})$

Knowing that our goal is to compute  $p(x|\mathcal{D})$ , we do this by:

$$p(x|\mathcal{D}) = \int p(x, \theta|\mathcal{D})d\theta$$

then:

$$p(x|\mathcal{D}) = \int p(x|\theta)p(\theta|\mathcal{D})d\theta$$

## Bayesian estimation(2)

Suppose  $\mu$  is the only unknown parameter, and

$$p(x|\mu) \sim \mathcal{N}(\mu, \sigma^2).$$

Our prior knowledge we might have about  $\mu$  can be expressed by a known prior density  $p(\mu)$ .

$$p(\mu) \sim \mathcal{N}(\mu_0, \sigma_0^2)$$

To obtain the posterior probability:

$$\begin{aligned} p(\mu|\mathcal{D}) &= \frac{p(\mathcal{D}|\mu)p(\mu)}{\int p(\mathcal{D}|\mu)p(\mu)} \\ &= \alpha \prod_{k=1}^n p(x_k|\mu)p(\mu) \end{aligned}$$



## Bayesian estimation(3)

Since  $p(x_k|\mu) \sim \mathcal{N}(\mu, \sigma^2)$  and  $p(\mu) \sim \mathcal{N}(\mu_0, \sigma_0^2)$ :

$$\begin{aligned}
 p(\mu|\mathcal{D}) &= \alpha \prod_{k=1}^n \overbrace{\frac{1}{\sqrt{2\pi}\sigma} \exp \left[ -\frac{1}{2} \left( \frac{x_k - \mu}{\sigma} \right)^2 \right]}^{p(x_k|\mu)} \overbrace{\frac{1}{\sqrt{2\pi}\sigma_0} \exp \left[ -\frac{1}{2} \left( \frac{\mu - \mu_0}{\sigma_0} \right)^2 \right]}^{p(\mu)} \\
 &= \alpha' \exp \left[ -\frac{1}{2} \left( \sum_{k=1}^n \left( \frac{\mu - x_k}{\sigma} \right)^2 + \left( \frac{\mu - \mu_0}{\sigma_0} \right)^2 \right) \right] \\
 &= \alpha'' \exp \left[ -\frac{1}{2} \left[ \left( \frac{n}{\sigma^2} + \frac{1}{\sigma_0^2} \right) \mu^2 - 2 \left( \frac{1}{\sigma^2} \sum_{k=1}^n x_k + \frac{\mu_0}{\sigma_0^2} \right) \mu \right] \right], \quad (
 \end{aligned}$$

Now we know  $p(\mu|\mathcal{D})$  is a Gaussian distribution, If we write  $p(\mu|\mathcal{D} \sim \mathcal{N}(\mu_n, \sigma_n^2))$ , we obtain:

$$\mu_n = \left( \frac{n\sigma_0^2}{n\sigma_0^2 + \sigma^2} \right) \bar{x}_n + \frac{\sigma^2}{n\sigma_0^2 + \sigma^2} \mu_0$$

$$\sigma_n^2 = \frac{\sigma_0^2 \sigma^2}{n\sigma_0^2 + \sigma^2}$$

Now we need to compute the  $p(x|\mathcal{D})$ :

$$\begin{aligned} p(x|\mathcal{D}) &= \int p(x|\mu)p(\mu|\mathcal{D}) d\mu \\ &= \int \frac{1}{\sqrt{2\pi}\sigma} \exp \left[ -\frac{1}{2} \left( \frac{x-\mu}{\sigma} \right)^2 \right] \frac{1}{\sqrt{2\pi}\sigma_n} \exp \left[ -\frac{1}{2} \left( \frac{\mu-\mu_n}{\sigma_n} \right)^2 \right] d\mu \\ &= \frac{1}{2\pi\sigma\sigma_n} \exp \left[ -\frac{1}{2} \frac{(x-\mu_n)^2}{\sigma^2 + \sigma_n^2} \right] f(\sigma, \sigma_n), \end{aligned} \quad (1)$$

Hence,

$$p(x|\mathcal{D}) \sim \mathcal{N}(\mu_n, \sigma^2 + \sigma_n^2)$$

# Margin(1)

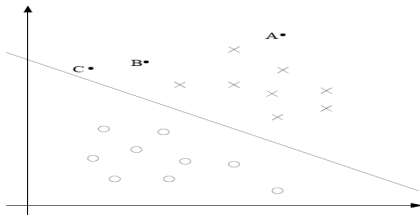
Previous:

Consider logistic regression, we compute  $w^T x$  then we:

- Predict "1" iff  $w^T x \geq 0$ .
- Predict "0" iff  $w^T x < 0$ .

The larger  $w^T x$  is, the larger also is  $h_w(x) = p(y|x; w)$ :

- if  $w^T x \gg 0$  then we are very "confident" predicting "1".
- if  $w^T x \ll 0$  then we are very "confident" predicting "0".



## Margin(2)

To make our discussion of SVMs easier, we need to change our notation:

we use  $y \in \{1, -1\}$  to denote class label, and change classifier as:

$$h_{w,b}(x) = g(w^T x + b)$$

Here  $g$  is a threshold function.

Given a training example  $(x^{(i)}, y^{(i)})$ , we define **functional margin** of  $(w, b)$  with respect to the training example.

$$\hat{\gamma}^{(i)} = y^{(i)}(w^T x^{(i)} + b)$$

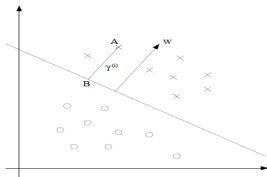
A large functional margin represents a confident and a correct prediction. But by exploiting our freedom to scale  $w$  and  $b$ , we can make the functional margin arbitrarily large without really changing anything meaningful.

Given a training set  $S$ , we also define the function margin of  $(w, b)$  with respect to  $S$ :

$$\hat{\gamma} = \min_{i=1,\dots,m} \hat{\gamma}^{(i)}$$

## Margin(3)

Lets talk about **geometric margins**. Consider the picture below:



Suppose A is  $x^{(i)}$ , then B is  $x^{(i)} - \gamma^{(i)} \cdot (w/\|w\|)$ , B lies on the decision boundary, Hence:

$$w^T \left( x^{(i)} - \gamma^{(i)} \cdot \frac{w}{\|w\|} \right) + b = 0$$

Solving for  $\gamma^{(i)}$  yields:

$$\gamma^{(i)} = \left( \frac{w}{\|w\|} \right)^T x^{(i)} + \frac{b}{\|w\|}$$

## Margin(4)

More generally, define geometric margins of  $(w, b)$  with respect to a training example  $(x^{(i)}, y^{(i)})$  to be

$$\gamma^{(i)} = y^{(i)} \left( \left( \frac{w}{\|w\|} \right)^T x^{(i)} + \frac{b}{\|w\|} \right)$$

Given a training set  $S$ , we also define the geometric margin of  $(w, b)$  with respect to  $S$  to be the smallest of the geometric margins on the individual training examples:

$$\gamma = \min_{i=1, \dots, m} \gamma^{(i)}$$

# The optimal margin classifier

Given a training set, We try to find a decision boundary that maximizes the margin. We can pose the following optimization problem

$$\begin{aligned} \max_{\gamma, w, b} \quad & \gamma \\ \text{s.t.} \quad & y^{(i)}(w^T x^{(i)} + b) \geq \gamma, i = 1, \dots, m \\ & \|w\| = 1. \end{aligned}$$

The " $\|w\|$ " constraint is a nasty(non-convex) constraint.  
Transform to a nicer one:

$$\begin{aligned} \max_{\gamma, w, b} \quad & \frac{\hat{\gamma}}{\|w\|} \\ \text{s.t.} \quad & y^{(i)}(w^T x^{(i)} + b) \geq \hat{\gamma}, i = 1, \dots, m \end{aligned}$$

We get rid of " $\|w\| = 1$ " which we don't like. But the downside is that we now have a nasty(non-convex) objective function.

## The optimal margin classifier(2)

Since we can add an arbitrary scaling constraint on  $w$  and  $b$  without changing anything. We introduce a scaling constraint:

$$\hat{\gamma} = 1$$

Plugging this into our problem above, and noting that maximizing  $\hat{\gamma}/\|w\| = 1/\|w\|$  is the same thing as minimizing  $\|w\|^2$ , we now have the following optimization problem:

$$\begin{aligned} \min_{\gamma, w, b} \quad & \frac{1}{2} \|w\|^2 \\ \text{s.t.} \quad & y^{(i)}(w^T x^{(i)} + b) \geq 1, i = 1, \dots, m \end{aligned}$$

Now, we have an optimization problem with a convex quadratic objective and only linear constraints. This optimization problem can be solved using commercial quadratic programming (QP) code.



# Lagrange duality

Consider problem of the following form:

$$\begin{array}{ll}\min_w & f(w) \\ s.t. & h_i(w) = 0, i = 1, \dots, l.\end{array}$$

Define **Lagrangian** to be

$$\mathcal{L}(w, \beta) = f(w) + \sum_{i=1}^l \beta_i h_i(w)$$

Here, the  $\beta_i$ 's are called the **Lagrange multipliers**. Set  $\mathcal{L}$ 's partial derivatives to zero:

$$\frac{\partial \mathcal{L}}{\partial w_i} = 0; \frac{\partial \mathcal{L}}{\partial \beta_i} = 0$$

We can solve for  $w$  and  $\beta$

## Lagrange duality(2)

Consider the following, which we call the **primal** optimization problem:

$$\begin{aligned} \min_w \quad & f(w) \\ \text{s.t.} \quad & g_i(w) \leq 0, i = 1, \dots, k \\ & h_i(w) = 0, i = 1, \dots, l \end{aligned}$$

Define **generalized Lagrangian**

$$\mathcal{L}(w, \alpha, \beta) = f(w) + \sum_{i=1}^k \alpha_i g_i(w) + \sum_{i=1}^l \beta_i h_i(w).$$

Here the  $\alpha_i$ 's and  $\beta_i$ 's are the Lagrange multipliers. Consider the quantity:

$$\theta_{\mathcal{P}} = \max_{\alpha, \beta: \alpha_i \geq 0} f(w) + \sum_{i=1}^k \alpha_i g_i(w) + \sum_{i=1}^l \beta_i h_i(w)$$

Here, the " $\mathcal{P}$ " subscript stands for "primal".

## Lagrange duality(3)

Let some  $w$  be given. If  $w$  violates any of the primal constraints, then it's easy to verify that  $\theta_{\mathcal{P}} = \infty$ . Conversely, if the constraints are indeed satisfied for a particular value of  $w$ , then  $\theta_{\mathcal{P}} = f(w)$ . Hence,

$$\theta_{\mathcal{P}}(w) = \begin{cases} f(w) & \text{if } w \text{ satisfied primal constraints} \\ \infty & \text{otherwise} \end{cases}$$

If we consider the minimization problem

$$\min_w \theta_{\mathcal{P}}(w) = \min_w \max_{\alpha, \beta, \alpha_i \geq 0} \mathcal{L}(w, \alpha, \beta)$$

It is the same problem of primal problem. Define optimal value of the objective to be  $p^* = \min_w \theta_{\mathcal{P}}(w)$

## Lagrange duality(4)

Define:

$$\theta_{\mathcal{D}}(\alpha, \beta) = \min_w \mathcal{L}(w, \alpha, \beta).$$

Here, the " $\mathcal{D}$ " subscript stands for "dual".

Now the **dual** optimization problem:

$$\max_{\alpha, \beta: \alpha_i \geq 0} \theta_{\mathcal{D}}(\alpha, \beta) = \max_{\alpha, \beta: \alpha_i \geq 0} \min_w \mathcal{L}(w, \alpha, \beta)$$

define the optimal value of the dual problems objective to be

$$d^* = \max_{\alpha, \beta: \alpha_i \geq 0} \theta_{\mathcal{D}}(\alpha, \beta).$$

It can easily be shown that

$$d^* = \max_{\alpha, \beta: \alpha_i \geq 0} \min_w \mathcal{L}(w, \alpha, \beta) \leq \min_w \max_{\alpha, \beta: \alpha_i \geq 0} \mathcal{L}(w, \alpha, \beta) = p^*$$

However under certain conditions, we will have  $d^* = p^*$

## Lagrange duality(5)

Let's see what these conditions are:

- $f$  and the  $g_i$ 's are convex;
- $h_i$ 's are affine;
- $g_i$ 's are (strictly) feasible; this means that there exists some  $w$  so that  $g_i(w) < 0$  for all  $i$ .

$w^*, \alpha^*$  and  $\beta^*$  satisfy the **Karush-Kuhn-Tucker (KKT) conditions**, which are as follows:

$$\frac{\partial}{\partial w_i} \mathcal{L}(w^*, \alpha^*, \beta^*) = 0, i = 1, \dots, n$$

$$\frac{\partial}{\partial \beta_i} \mathcal{L}(w^*, \alpha^*, \beta^*) = 0, i = 1, \dots, l$$

$$\alpha_i^* g_i(w^*) = 0, i = 1, \dots, k$$

$$g_i(w^*) \leq 0, i = 1, \dots, k$$

$$\alpha_i^* \geq 0, i = 1, \dots, k$$

## The optimal margin classifier(3)

we posed the following (primal) optimization problem for finding the optimal margin classifier:

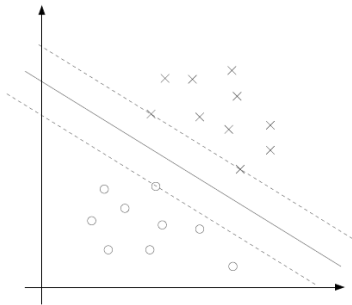
$$\begin{aligned} \min_{\gamma, w, \beta} \quad & \frac{1}{2} \|w\|^2 \\ \text{s.t.} \quad & y^{(i)}(w^T x^{(i)} + b) \geq 1, i = 1, \dots, m \end{aligned}$$

We can write the constraints as:

$$g_i(x) = 1 - y^{(i)}(w^T x^{(i)} + 1) \leq 0, i = 1, \dots, m$$

Note that from the KKT condition, we will have  $\alpha_i > 0$  only for that training example that have function margin exactly equal to one

## The optimal margin classifier(4)



The points with the smallest margins are exactly the ones closest to the decision boundary; here, these are the three points which are called the **support vectors** in this problem.

## The optimal margin classifier(5)

Construct Lagrangian for our problem:

$$\mathcal{L}(w, b, \alpha) = \frac{1}{2} \|w\|^2 - \sum_{i=1}^m [y^{(i)}(w^T x^{(i)} + b) - 1]$$

Note that there are only  $\alpha_i$  but no  $\beta_i$  Lagrange multipliers. To minimize  $\mathcal{L}(w, b, \alpha)$  for fix  $\alpha$ :

$$\nabla_w \mathcal{L}(w, b, \alpha) = w - \sum_{i=1}^m \alpha_i y^{(i)} x^{(i)} \stackrel{\text{set}}{=} 0$$

This implies that:

$$w = \sum_{i=1}^m \alpha_i y^{(i)} x^{(i)} \tag{1}$$



## The optimal margin classifier(6)

As for the derivation with respect to  $b$ , we obtain:

$$\frac{\partial}{\partial b} \mathcal{L}(w, b, \alpha) = \sum_{i=1}^m \alpha_i y^{(i)} = 0 \quad (2)$$

Plugging  $w$  back into the Lagrangian:

$$\mathcal{L}(w, b, \alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j \langle x^{(i)}, x^{(j)} \rangle$$

We obtain the following dual problem:

$$\begin{aligned} \max_{\alpha} \quad & W(\alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j \langle x^{(i)}, x^{(j)} \rangle \\ \text{s.t.} \quad & \alpha_i \geq 0, \quad i = 1, \dots, m \\ & \sum_{i=1}^m \alpha_i y^{(i)} = 0 \end{aligned}$$

## The optimal margin classifier(7)

Having found  $w^*$ , then we can get  $b^*$  as:

$$b^* = - \frac{\max_{i:y^{(i)}=-1} w^{*T} x^{(i)} + \min_{i:y^{(i)}=1} w^{*T} x^{(i)}}{2}$$

Suppose we've fit our model's parameters to a training set, and now wish to make a prediction at a new point  $x$ . We would calculate  $w^T x + b$ , and do prediction. But using (1), this quantity can also be written:

$$w^T x + b = \sum_{i=1}^m \alpha_i y^{(i)} \langle x^{(i)}, x \rangle + b$$

Using (2) we saw that the  $\alpha_i$ 's will all be zero except for the support vector.

# Kernels

Define **Kernel** to be:

$$K(x, z) = \phi(x)^T \phi(z)$$

we let  $\phi$  denote the **feature mapping**, which maps from the original input value (**attributes**) to the features in higher dimensional.

Let's see an example. Suppose  $x, z \in \mathbb{R}^n$ , and consider:

$$K(x, z) = (x^T z)^2 = \sum_{i,j=1}^n (x_i x_j)(z_i z_j) = \phi(x)^T \phi(z)$$

Thus, we can see feature mapping  $\phi$  is given by (shown here for the case of  $n = 3$ )

$$\phi(x) = (x_1 x_1, x_1 x_2, \dots, x_3 x_3)^T$$

## Kernels(2)

Another example:

$$\begin{aligned} K(x, z) &= (x^T z + c)^2 \\ &= \sum_{i,j=1}^n (x_i x_j)(z_i z_j) + \sum_{i=1}^n (\sqrt{2c} x_i)(\sqrt{2c} z_i) + c^2 \end{aligned}$$

feature mapping: (Again shown for  $n = 3$ )

$$\phi(x) = (x_1 x_1, x_1 x_2, \dots, x_3 x_3, \sqrt{2c} x_1, \sqrt{2c} x_2, \sqrt{2c} x_3, c)^T$$

More broadly, the kernel  $K(x, z) = (x^T z + c)^d$  corresponds to a feature mapping to an  $\binom{n+d}{d}$  feature space.

Map attribute to infinity dimension **Gaussian kernel**:

$$K(x, z) = \exp\left(-\frac{\|x - z\|^2}{2\sigma^2}\right)$$

## Kernels(3)

More broadly, given some function  $K$ , how can we tell if it's a valid kernel; i.e., can we tell if there is some feature mapping  $\phi$  so that  $K(x, z) = \phi(x)^T \phi(z)$  for all  $x, z$ ?

Consider some finite set of  $m$  points  $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$ , and let a m-by-m matrix  $K$  be defined as

$$K_{ij} = K(x^{(i)}, y(i))$$

$K$  is called the **Kernel matrix**.

Now, if  $K$  is a valid Kernel, then

$$K_{ij} = \phi(x^{(i)})^T \phi(x^{(j)}) = \phi(x^{(j)})^T \phi(x^{(i)}) = K_{ji}$$

Hence,  $K$  must be symmetric.

## Kernels(4)

More over, letting  $\phi_k(x)$  denote the  $k$ -th coordinate of the vector  $\phi(x)$ , we find that for any vector  $z$ , we have

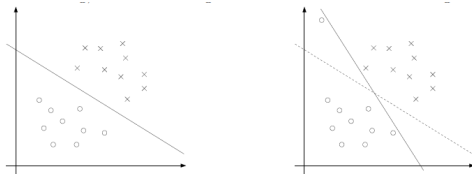
$$\begin{aligned} z^T K z &= \sum_i \sum_j z_i K_{ij} z_j \\ &= \sum_i \sum_j z_i \phi(x^{(i)})^T \phi(x^{(j)}) z_j \\ &= \sum_i \sum_j z_i \sum_k \phi_k(x^{(i)})^T \phi_k(x^{(j)}) z_j \\ &= \sum_k \sum_i \sum_j z_i \phi_k(x^{(i)})^T \phi_k(x^{(j)}) z_j \\ &= \sum_k \left( \sum_i z_i \phi_k(x^{(i)}) \right)^2 \geq 0 \end{aligned}$$

This shows that  $K$  is positive semi-definite

## Kernels(5)

**Theorem (Mercer).** Let  $K : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$  be given. Then for  $K$  to be a valid (Mercer) Kernel, it is necessary and sufficient that for any  $\{x^{(1)}, \dots, x^{(m)}\}$ ,  $(m < \infty)$ , the corresponding kernel matrix is symmetric positive semi-definite.

## Regularization and the non-separable case



To make the algorithm work for non-linearly separable datasets as well as be less sensitive to outliers, we reformulate our optimization ( $l_1$  **regularization**):

$$\begin{aligned} \min_{\gamma, w, b} \quad & \frac{1}{2} \|w\|^2 + C \sum_{i=1}^m \xi_i \\ \text{s.t.} \quad & y^{(i)}(w^T x^{(i)} + b) \geq 1 - \xi_i, i = 1, \dots, m \\ & \xi_i \geq 0, i = 1, \dots, m \end{aligned}$$



## Regularization and the non-separable case(2)

As before, we can form the Lagrangian:

$$\mathcal{L}(w, b, \xi, \alpha, \gamma) = \frac{1}{2}w^T w + C \sum_{i=1}^m \xi_i - \sum_{i=1}^m \alpha_i (y^{(i)}(x^T w + b) - 1 + \xi_i) - \sum_{i=1}^m \gamma_i \xi_i$$

Here, the  $\alpha_i$ 's and  $\gamma_i$ 's are our Lagrangian multipliers.

After some work, we obtain the dual form of the problem:

$$\begin{aligned} \max_{\alpha} \quad & W(\alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j \langle x^{(i)}, x^{(j)} \rangle \\ \text{s.t.} \quad & 0 \leq \alpha_i \leq C, i = 1, \dots, m \\ & \sum_{i=1}^m \alpha_i y^{(i)} = 0 \end{aligned}$$

# The SMO algorithm

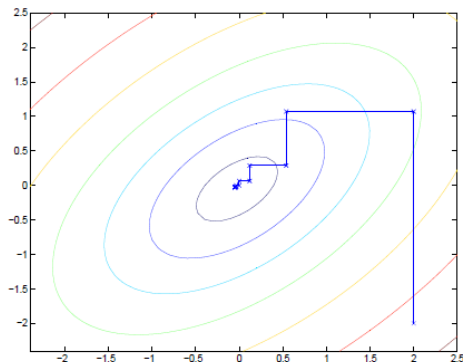
Consider trying to solve the unconstrained optimization problem

$$\max_{\alpha} W(\alpha_1, \alpha_2, \dots, \alpha_m).$$

The algorithm here is called **coordinate ascent**:

```
Loop until convergence: {  
  For  $i = 1, \dots, m$ , {  
     $\alpha_i := \arg \max_{\hat{\alpha}_i} W(\alpha_1, \dots, \alpha_{i-1}, \hat{\alpha}_i, \alpha_{i+1}, \dots, \alpha_m).$   
  }  
}
```

## The SMO algorithm(2)



The ellipses in the figure are the contours of a quadratic function that we want to optimize.

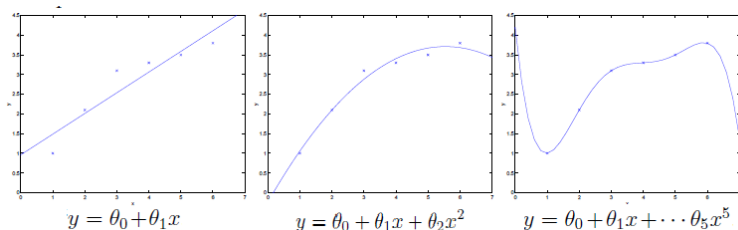
# The SMO algorithm(3)

Repeat till convergence {

1. Select some pair  $\alpha_i$  and  $\alpha_j$  to update next (using a heuristic that tries to pick the two that will allow us to make the biggest progress towards the global maximum).
2. Reoptimize  $W(\alpha)$  with respect to  $\alpha_i$  and  $\alpha_j$ , while holding all the other  $\alpha_k$ 's ( $k \neq i, j$ ) fixed.

}

# Bias/variance tradeoff



- The linear model suffers from large bias, and may underfit the data.
- The 5th order polynomial model suffers from large variance, and may overfit the data.

## Preliminaries

**Lemma.** (The union bound). Let  $A_1, A_2, \dots, A_k$  be  $k$  different events (that may not be independent). Then

$$P(A_1 \cup \dots \cup A_k) \leq P(A_1) + \dots + P(A_k)$$

**Lemma.** (Hoeffding inequality) Let  $Z_1, \dots, Z_m$  be the  $m$  independent and identically distributed (iid) random variables drawn from a Bernoulli( $\phi$ ) distribution. I.e.,  $P(Z_i = 1) = \phi$ , and  $P(Z_i = 0) = 1 - \phi$ . Let  $\hat{\phi} = (1/m) \sum_{i=1}^m Z_i$  be the mean of these random variables, and let any  $\gamma > 0$  be fixed. Then

$$P(|\phi - \hat{\phi}| > \gamma) \leq 2 \exp(-2\gamma^2 m)$$

## Preliminaries(2)

We assume we are given a training set

$S = \{(x^{(i)}, y^{(i)}); i = 1, \dots, m\}$  of size  $m$ , where the training examples  $(x^{(i)}, y^{(i)})$  are drawn iid from some probability distribution  $\mathcal{D}$ . For a hypothesis  $h$ , we define the **training error** (also called the **empirical risk** or **empirical error** in learning theory) to be

$$\hat{\varepsilon}(h) = \frac{1}{m} \sum_{i=1}^m 1\{h(x^{(i)}) \neq y^{(i)}\}$$

Here  $1\{condition\}$  means, if condition is true then  $1\{condition\} = 1$ , otherwise  $1\{condition\} = 0$ .

Define **generalization error** to be

$$\varepsilon(h) = P_{(x,y) \sim \mathcal{D}}(h(x) \neq y)$$

# Empirical risk minimization (ERM)

We define the **hypothesis class**  $\mathcal{H}$  used by a learning algorithm to be the set of all classifiers considered by it. For linear classification,  $\mathcal{H} = \{h_w : h_w(x) = 1\{w^T x \geq 0\}, w \in \mathbb{R}^{n+1}\}$ .

Empirical risk minimization can now be thought of as a minimization over the class of functions  $\mathcal{H}$ :

$$\hat{h} = \arg \min_{h \in \mathcal{H}} \hat{\varepsilon}(h)$$

Let's consider the case of finite  $\mathcal{H} = \{h_1, \dots, h_k\}$ . Take any one, fixed,  $h_i \in \mathcal{H}$ . Consider a Bernoulli random variable  $Z_j = 1\{h_i(x^{(j)}) \neq y^{(j)}\}$ , then the training error can be written

$$\hat{\varepsilon}(h_i) = \frac{1}{m} \sum_{j=1}^m Z_j$$

Since our training set was drawn iid from  $\mathcal{D}$

$$\varepsilon(h_i) = E[Z_j] = P(Z_j = 1)$$



## Empirical risk minimization (ERM)(2)

We can apply the Hoeffding inequality, and obtain

$$P(|\varepsilon(h_i) - \hat{\varepsilon}(h_i)| > \gamma) \leq 2 \exp(-2\gamma^2 m)$$

Let  $A_i$  denote the event that  $|\varepsilon(h_i) - \varepsilon(\hat{h}_i)| > \gamma$ . Using the union bound, we have

$$\begin{aligned} P(\exists h \in \mathcal{H}. |\varepsilon(h_i) - \hat{\varepsilon}(h_i)| > \gamma) &= P(A_1 \cup \dots \cup A_k) \\ &\leq \sum_{i=1}^k P(A_i) \\ &\leq \sum_{i=1}^k 2 \exp(-2\gamma^2 m) \\ &= 2k \exp(-2\gamma^2 m) \end{aligned}$$

## Empirical risk minimization (ERM)(3)

Subtract both sides from 1, we obtain

$$\begin{aligned} P(\neg \exists h \in \mathcal{H}. |\varepsilon(h_i) - \hat{\varepsilon}(h_i)| > \gamma) &= P(\forall h \in \mathcal{H}. |\varepsilon(h_i) - \hat{\varepsilon}(h_i)| \leq \gamma) \\ &\geq 1 - 2k \exp(-2\gamma^2 m) \end{aligned}$$

So, with probability at least  $1 - 2k \exp(-2\gamma^2 m)$  we have that  $\varepsilon(h)$  will be within  $\gamma$  of  $\varepsilon(\hat{h})$  for all  $h \in \mathcal{H}$ . This is called a **uniform convergence result**.

Given  $\gamma$  and some  $\delta > 0$ , by setting  $\delta = 2k \exp(-2\gamma^2 m)$ , we find that if

$$m \geq \frac{1}{2\gamma^2} \log \frac{2k}{\delta},$$

then with probability at least  $1 - \delta$ , we have that  $|\varepsilon(h) - \hat{\varepsilon}(h)| \leq \gamma$  for all  $h \in \mathcal{H}$ .

The training set size  $m$  that a certain method or algorithm requires in order to achieve a certain level of performance is also called the algorithm's **sample complexity**.

## Empirical risk minimization (ERM)(4)

Similarly, we can also hold  $m$  and  $\delta$  fixed and solve for  $\gamma$ .

$$|\varepsilon(\hat{h}) - \varepsilon(h)| \leq \gamma \leq \sqrt{\frac{1}{2m} \log \frac{2k}{\delta}}$$

Define  $h^* = \arg \min_{h \in \mathcal{H}} \varepsilon(h)$ . We have:

$$\varepsilon(\hat{h}) \leq \hat{\varepsilon}(\hat{h}) + \gamma \leq \hat{\varepsilon}(h^*) + \gamma \leq \varepsilon(h^*) + 2\gamma$$

**Theorem.** Let  $|\mathcal{H}| = k$ , and let any  $m, \delta$  be fixed. Then with probability at least  $1 - \delta$ , we have that

$$\varepsilon(\hat{h}) \leq \left( \min_{h \in \mathcal{H}} \varepsilon(h) \right) + 2\sqrt{\frac{1}{2m} \log \frac{2k}{\delta}}$$

## Empirical risk minimization (ERM)(5)

Given a set  $S = \{x^{(1)}, \dots, x^{(d)}\}$ , we say that  $\mathcal{H}$  **shatters**  $S$  if  $\mathcal{H}$  can realize any labelling on  $S$ . I.e., if for any set of labels  $\{y^{(1)}, \dots, y^{(d)}\}$ , there exists some  $h \in \mathcal{H}$  so that  $h(x^{(i)}) = y^{(i)}$  for all  $i = 1, \dots, d$ .

Given a hypothesis class  $\mathcal{H}$ , we can define its **Vapnik-Chervonenkis dimension**, written  $VC(\mathcal{H})$ , to be the size of the largest set that is shattered by  $\mathcal{H}$ .

**Theorem.** Let  $\mathcal{H}$  be given, and let  $d = VC(\mathcal{H})$ . Then with probability at least  $1 - \delta$ , we have that for all  $h \in \mathcal{H}$ ,

$$|\epsilon(h) - \hat{\epsilon}(h)| \leq O \left( \sqrt{\frac{d}{m} \log \frac{m}{d} + \frac{1}{m} \log \frac{1}{\delta}} \right).$$

## Empirical risk minimization (ERM)(6)

Thus, with probability at least  $1 - \delta$ , we also have that:

$$\epsilon(\hat{h}) \leq \epsilon(h^*) + O\left(\sqrt{\frac{d}{m} \log \frac{m}{d} + \frac{1}{m} \log \frac{1}{\delta}}\right).$$

**Corollary.** For  $|\epsilon(h) - \hat{\epsilon}(h)| \leq \gamma$  to hold all  $h \in \mathcal{H}$  with probability at least  $1 - \delta$ , it suffices that  $m = O_{\gamma,\delta}(d)$ .

In Mixture densities we have some assumptions.

1. The samples come from a known number of classes.
2. The prior probabilities  $P(\omega_j)$  for each class are known,  $j = 1, \dots, c$ .
3. The forms for the class-conditional probability densities  $p(x|\omega_j, \theta_j)$  are known,  $j = 1, \dots, c$ .
4. The values for the  $c$  parameter vectors.  $\theta_1, \dots, \theta_c$  are unknown.
5. The category labels are unknown.

Samples are assumed to be obtained by selecting state if nature  $\omega_j$  with probability  $P(\omega_j)$  and then selecting an  $x$  according to the probability law  $p(x|\omega_j, \theta_j)$ . Thus probability density function for the samples is given by

$$p(x|\theta) = \sum_{j=1}^c p(x|\omega_j, \theta_j)P(\omega_j)$$

**Definition:** A density  $p(x|\theta)$  is said to be *identifiable* if  $\theta \neq \theta'$  implies that there exists an  $x$  such that  $p(x|\theta) \neq p(x|\theta')$ .

Or put another way, a density  $p(x|\theta)$  is *not identifiable* if we cannot recover a unique  $\theta$ .

A example for unidentifiable:

$$\begin{aligned} P(x|\theta) &= \frac{1}{2}\theta_1^x(1-\theta_1)^{1-x} + \frac{1}{2}\theta_2^x(1-\theta_2)^{1-x} \\ &= \begin{cases} \frac{1}{2}(\theta_1 + \theta_2) & \text{if } x = 1; \\ 1 - \frac{1}{2}(\theta_1 + \theta_2) & \text{if } x = 0. \end{cases} \end{aligned}$$

Suppose, for example, that we know for our data that  $P(x = 1) = 0.6$ , and hence that  $P(x = 0|\theta) = 0.4$ . Then we know the function  $P(x|\theta)$ , but we cannot determine  $\theta$ , and hence cannot extract the component distributions. The most we can say is that  $\theta_1 + \theta_2 = 1.2$ . Thus, here we have a case in which the mixture distribution is completely unidentifiable.

Suppose now that we are given a set  $\mathcal{D} = \{x_1, \dots, x_n\}$  of  $n$  unlabeled samples drawn independently from mixture density.

$$p(x|\theta) = \sum_{j=1}^c p(x|\omega_j, \theta_j) P(\omega_j)$$

where the full parameter vector  $\theta$  is fixed but unknown. The likelihood of the observed samples is, by definition, the joint density

$$p(\mathcal{D}|\theta) \equiv \prod_{k=1}^n p(x_k|\theta)$$

The maximum-likelihood estimate  $\hat{\theta}$  is that value of  $\theta$  that maximizes  $p(\mathcal{D}|\theta)$ .



The Log maximum-likelihood:

$$l = \sum_{k=1}^n \log p(x_k | \theta)$$

and

$$\begin{aligned} \nabla_{\theta_i} l &= \sum_{k=1}^n \frac{1}{p(x_k | \theta)} \nabla_{\theta_i} \left( \sum_{j=1}^c p(x_k | \omega_j, \theta_j) P(\omega_j) \right) \\ &= \sum_{k=1}^n \frac{1}{p(x_k | \theta)} \nabla_{\theta_i} p(x_k | \omega_i) P(\omega_i) \end{aligned}$$

Here we assume that  $\theta_i$  and  $\theta_j$  are independent if  $i \neq j$ .

According the Bayes' formula,

$$P(\omega_i|x_k, \theta) = \frac{p(x_k|\omega_i, \theta_i)P(\omega_i)}{p(x_k|\theta)}$$

Hence,

$$\nabla_{\theta_i} l = \sum_{k=1}^n P(\omega_i|x_k, \theta) \nabla_{\theta_i} \log p(x_k|\omega_i, \theta_i)$$

and  $P(\omega_i)$  can be estimate by  $P(\omega_i|x_k)$ :

$$p(\omega_i) = \frac{1}{n} \sum_{k=1}^n P(\omega_i|x_k, \theta)$$

# Mixtures of Gaussians and the EM algorithm

Suppose that we are given a training set  $\{x^{(1)}, \dots, x^{(m)}\}$ , we wish to model the data by specifying a joint distribution  $p(x^{(1)}, z^{(i)}) = p(x^{(i)}|z^{(i)})$ . Here,

$$\begin{aligned} z^{(i)} \in \{1, \dots, k\} &\sim \text{Multinomial}(\phi), \\ x^{(i)}|z^{(i)} = j &\sim \mathcal{N}(\mu_j, \Sigma_j) \end{aligned}$$

This is call the **mixture of Gaussians** model, and  $z^{(i)}$ 's are **latent** random variables. To estimate  $\phi$ ,  $\mu$  and  $\Sigma$ , we can write down the likelihood of our data.

$$\begin{aligned} l(\phi, \mu, \Sigma) &= \sum_{i=1}^m \log p(x^{(i)}; \phi, \mu, \Sigma) \\ &= \sum_{i=1}^m \log \sum_{z^{(i)}=1}^k p(x^{(i)}|z^{(i)}; \mu, \Sigma) p(z^{(i)}; \phi) \end{aligned}$$

## Mixtures of Gaussians and the EM algorithm(2)

Note that if we knew what the  $z^{(i)}$ 's were, the maximum likelihood problem would have been easy. Specifically, we could then write down the likelihood as

$$l(\phi, \mu, \Sigma) = \sum_{i=1}^m \log p(x^{(i)} | z^{(i)}; \mu, \Sigma) + \log p(z^{(i)}; \phi)$$

Maximizing this with respect to  $\phi$ ,  $\mu$  and  $\Sigma$  gives the parameters.

$$\begin{aligned}\phi_j &= \frac{1}{m} \sum_{i=1}^m 1\{z^{(i)} = j\} \\ \mu_j &= \frac{\sum_{i=1}^m 1\{z^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m 1\{z^{(i)} = j\}} \\ \Sigma_j &= \frac{\sum_{i=1}^m 1\{z^{(i)} = j\} (x^{(i)} - \mu_j)(x^{(i)} - \mu_j)^T}{\sum_{i=1}^m 1\{z^{(i)} = j\}}\end{aligned}$$

# Mixtures of Gaussians and the EM algorithm(3)

The EM algorithm is an iterative algorithm that has two main step.

Repeat until convergence: {

(E-step) For each  $i, j$ , set

$$w_j^{(i)} := p(z^{(i)} = j | x^{(i)}; \phi, \mu, \Sigma)$$

(M-step) Update the parameters:

$$\phi_j := \frac{1}{m} \sum_{i=1}^m w_j^{(i)},$$

$$\mu_j := \frac{\sum_{i=1}^m w_j^{(i)} x^{(i)}}{\sum_{i=1}^m w_j^{(i)}},$$

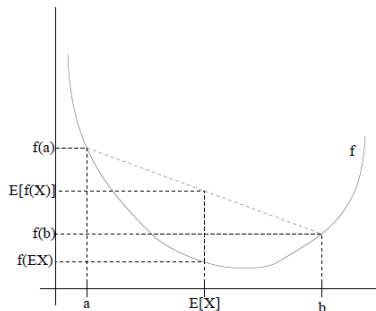
$$\Sigma_j := \frac{\sum_{i=1}^m w_j^{(i)} (x^{(i)} - \mu_j)(x^{(i)} - \mu_j)^T}{\sum_{i=1}^m w_j^{(i)}}$$

## Jensens inequality

**Theorem.** Let  $f$  be a convex(concave) function, and let  $X$  be a random variable. Then:

$$E[f(X)] \geq f(EX) \quad (\text{Concave case is } E[f(X)] \leq f(EX))$$

Moreover, if  $f$  is strictly convex(concave), then  $E[f(X)] = f(EX)$  holds true if and only if  $X = E[X]$  with probability 1.



# The EM algorithm

Suppose we have a training set  $\{x^{(i)}, \dots, x^{(m)}\}$ . We wish to fit the parameters of a model  $p(x, z)$  to the data, where the likelihood is given by

$$l(\theta) = \sum_{i=1}^m \log p(x; \theta) = \sum_{i=1}^m \log \sum_z p(x, z; \theta).$$

For each  $i$ , let  $Q_i$  be some distribution over the  $z$ 's ( $\sum_z Q_i(z) = 1, Q_i(z) \geq 0$ ). Consider:

$$\sum_i \log p(x^{(i)}; \theta) = \sum_i \log \sum_{z^{(i)}} p(x^{(i)}, z^{(i)}; \theta) \quad (3)$$

$$= \sum_i \log \sum_{z^{(i)}} Q_i(z^{(i)}) \frac{p(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})} \quad (4)$$

$$= \sum_i \log E_{z^{(i)} \sim Q_i} \left[ \frac{p(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})} \right] \quad (5)$$

$$\geq \sum_i \sum_{z^{(i)}} Q_i(z^{(i)}) \log \frac{p(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})} \quad (6)$$

## The EM algorithm(2)

To make the bound tight, we need Jensen's inequality to hold equality. we require that

$$\frac{p(x^{(i)}, z^{(i)})}{Q_i(z^{(i)})} = c$$

for some constant  $c$  that not depend on  $z^{(i)}$

$$Q_i(z^{(i)}) \propto p(x^{(i)}, z^{(i)}; \theta)$$

Actually, since  $\sum_z Q_i(z^{(i)}) = 1$ :

$$\begin{aligned} Q_i(z^{(i)}) &= \frac{p(x^{(i)}, z^{(i)}; \theta)}{\sum_z p(x^{(i)}, z; \theta)} \\ &= \frac{p(x^{(i)}, z^{(i)}; \theta)}{p(x^{(i)}; \theta)} = p(z^{(i)} | x^{(i)}; \theta) \end{aligned}$$



# The EM algorithm(3)

Generalized EM algorithm:

Repeat until convergence {

(E-step) For each  $i$ , set

$$Q_i(z^{(i)}) := p(z^{(i)}|x^{(i)}; \theta).$$

(M-step) Set

$$\theta := \arg \max_{\theta} \sum_i \sum_{z^{(i)}} Q_i(z^{(i)}) \log \frac{p(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})}.$$

}

If we define

$$J(Q, \theta) = \sum_i \sum_{z^{(i)}} Q_i(z^{(i)}) \log \frac{p(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})}.$$

the EM can also be viewed a coordinate ascent on  $J$ .

## The envelope quiz



- red ball is goal
- You randomly picked an envelop randomly took out a ball and it was black
- Should you choose this envelope or the other envelope?

## The envelope quiz(2)

### ■ Probabilistic inference

- Joint distribution on  $E \in \{1, 0\}, B \in \{r, b\}$ :  $P(E, B) = P(E)P(B|E)$
- $P(E = 1) = P(E = 0) = \frac{1}{2}$
- $P(B = r|E = 1) = \frac{1}{2}, P(B = r|E = 0) = 0$
- The graphical model:



- Statistical decision theory: switch if  $P(E = 1|B = b) < \frac{1}{2}$
- $P(E = 1|B = b) = \frac{P(B=b|E=1)P(E=1)}{P(B=b)} = \frac{1}{3}$

## Some concepts

- The world is reduced to a set of random variables  $x_1, \dots, x_d$ 
  - e.g.  $(x_1, \dots, x_{d-1})$  is a feature vector,  $x_d \equiv y$  is the class label.
- Inference: given joint distribution  $p(x_1, \dots, x_d)$ , compute  $p(X_Q|X_E)$  where  $X_Q \cup X_E \subseteq \{x_1, \dots, x_d\}$ 
  - e.g.  $Q = \{x_d\}, E = \{x_1, \dots, x_{d-1}\}$ , by the definition of conditional

$$p(x_d|x_1, \dots, x_{d-1}) = \frac{p(x_1, \dots, x_d)}{\sum_v p(x_1, \dots, x_{d-1}, x_d = v)}$$

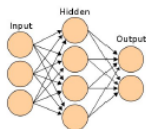
- Learning estimate  $p(x_1, \dots, x_d)$  from training data  $\{x^{(1)}, \dots, x^{(m)}\}$

# What are graphical models?

- Graphical model = joint distribution  $p(x_1, \dots, x_d)$ 
  - Bayesian network
  - Markov random field
- Inference =  $p(X_Q|X_E)$ , in general  $X_Q \cup X_E \subseteq \{x_1, \dots, x_d\}$
- If  $p(x_1, \dots, x_d)$  *not given, estimate it from data*
  - parameter and structure learning

## What are graphical models?(2)

- Graphical model is the study of probabilistic models
- Just because there are nodes and edges doesn't mean it's a graphical model
- These are not graphical models:



neural network



decision tree



network flow



HMM template  
(but HMMs are!)

## Directed graphical models

- Also called Bayesian networks
- A directed graph has nodes  $x_1, \dots, x_d$ , some of them connected by directed edges  $x_i \rightarrow x_j$
- A cycle is a directed path  $x_1 \rightarrow \dots \rightarrow x_k$  where  $x_1 = x_k$
- A directed acyclic graph (DAG) contains no cycles

## Directed graphical models(2)

- A Bayesian network on the DAG is a family of distributions satisfying

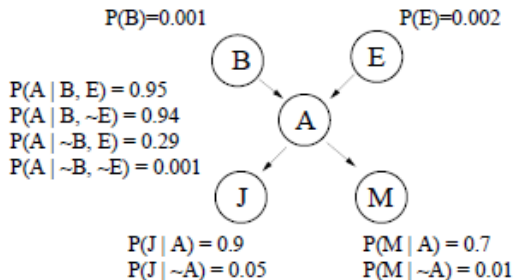
$$p|p(x_1, \dots, x_d) = \prod_i p(x_i | Pa(x_i))$$

where  $Pa(x_i)$  is the set of parents of  $x_i$

- $p(x_i | Pa(x_i))$  is the conditional probability distribution (CPD) at  $x_i$
- By specifying the CPDs for all  $i$ , we specify a joint distribution  $p(x_1, \dots, x_d)$

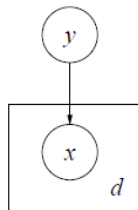
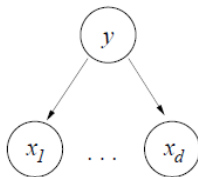


## Example: Burglary, Earthquake, Alarm, John and Marry



$$P(B, \sim E, A, J, \sim M) = P(B)P(\sim E)P(A|B, \sim E)P(J|A)P(\sim M|A)$$

## Example: Naive Bayes



- $p(y, x_1, \dots, x_d) = p(y) \prod_{i=1}^d p(x_i|y)$
- Plate representation on the right
- $p(y)$  multinomial
- $p(x_i|y)$  depends on the feature type.

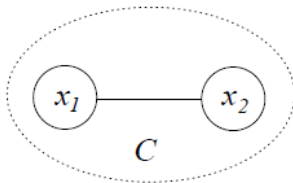
# Undirected graphical models

- Also known as Markov Random Fields
- A clique  $C$  in an undirected graph is a set of fully connected nodes (full of loops!)
- Define a nonnegative potential function  $\psi_C : X_C \rightarrow \mathbb{R}_+$
- An undirected graphical model is a family of distributions satisfying

$$\left\{ p | p(X) = \frac{1}{Z} \prod_C \psi_C(X_C) \right\}$$

- $Z = \int \prod_C \psi_C(X_C) dX$

# A Tiny Markov Random Field



- $x_1, x_2 \in \{-1, 1\}$
- A single clique  $\psi_C(x_1, x_2) = e^{ax_1x_2}$
- $p(x_1, x_2) = \frac{1}{Z} e^{ax_1x_2}$
- $Z = (2e^a + 2e^{-a})$
- $p(1, 1) = p(-1, -1) = \frac{e^a}{2e^a + 2e^{-a}}$

# Topic Model

- A topic is defined as a probability distribution over terms or a cluster of weighted terms.
- A document is defined as a set of words generated from a mixture of latent topics.

Various topic modeling methods, such as PLSI, LDA, LSI, NMF, and RLSI have been proposed and successfully applied to different applications.

# Probabilistic Topic Models

Suppose that  $\mathcal{D} = \{d_1, d_2, \dots, d_N\}$  is a set of documents with size  $N$  and  $\mathcal{V}$  is a set of terms or words with size  $M$ , i.e., the vocabulary. A document  $d \in \mathcal{D}$  consists of  $|d|$  words from the vocabulary, denoted as  $d = (w_1, w_2, \dots, w_{|d|})$ . Suppose that there are  $K$  topics in the document collection.

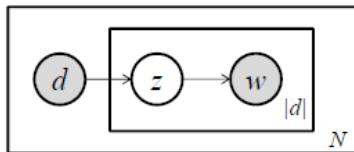
PLSI is one of the widely used probabilistic topic models. One can generate the documents in the collection in the following way.

- select a document  $d$  from the collection with probability  $P(d)$
- select a latent topic  $z$  with probability  $P(z|d)$
- generate a word  $w$  with probability  $P(w|z)$

Here  $z \in \{z_1, \dots, z_K\}$  is a latent variable representing a topic.

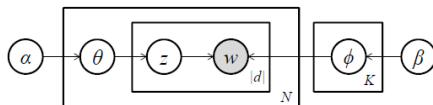
## Probabilistic Topic Models(2)

The parameters of  $P(d)$ ,  $P(w|z)$ , and  $P(z|d)$  can be estimated by EM algorithm.



# Probabilistic Topic Models(3)

LDA Model:



1. for each topic  $k = 1, \dots, K$ 
  - (a) draw word distribution  $\phi_k$  according to  $\phi_k | \beta \sim \text{Dir}(\beta)$
2. for each document  $d$  in the collection
  - (a) draw topic distribution  $\theta$  according to  $\theta | \alpha \sim \text{Dir}(\alpha)$
  - (b) for each word  $w$  in the document  $d$ 
    - i. draw a topic  $z$  according to  $z | \theta \sim \text{Mult}(\theta)$
    - ii. draw a word  $w$  according to  $w | z \sim \text{Mult}(\phi_z)$



## Non-probabilistic Topic Models

Non-probabilistic topic models are usually obtained by matrix factorization. Suppose that  $\mathcal{D}$  is a set of document with size  $N$ , and  $\mathcal{V}$  is a vocabulary with size  $M$ . The document collection  $\mathcal{D}$  is represented as an  $M \times N$  matrix  $D = (d_1, \dots, d_N)$ .

Suppose that there are  $K$  topics represented as an  $M \times K$  term-topic matrix  $U = (u_1, \dots, u_K)$ .  $V^T = (v_1, \dots, v_N)$  is an  $K \times N$  topic-document matrix.

Our goal is to factor matrix  $D$  by  $U$  and  $V^T$ .

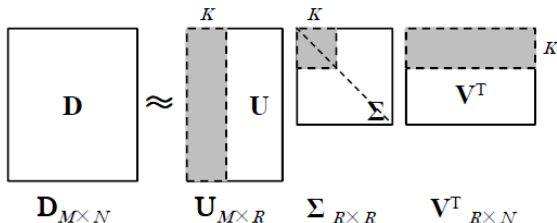
$$D \approx UV^T$$

## Non-probabilistic Topic Models(2)

LSI assumes that the  $K$  columns of matrix  $U$  as well as the  $K$  columns of matrix  $V$  are orthonormal. LSI amounts to minimizing the following objective function with the orthonormal constraints.

$$\min_{U, V} \|D - U\Sigma V^T\|_F$$

*s.t.*  $U^T \times U = I, V \times V^T = I$ , and  $\Sigma$  is diagonal



## What is entropy

In thermodynamics, entropy (usual symbol  $S$ ) is a measure of the number of specific ways in which a thermodynamic system may be arranged, commonly understood as a measure of disorder.

$$\Delta = \int \frac{dQ_{rev}}{T}$$

In information theory, **Entropy** is a measure of unpredictability of information content. Shannon defined the entropy  $H$  of a discrete random variable  $X$  with possible values  $x_1, \dots, x_n$  and probability mass function  $P(X)$  as:

$$H(X) = E[I(X)] = E[-\ln(P(X))]$$

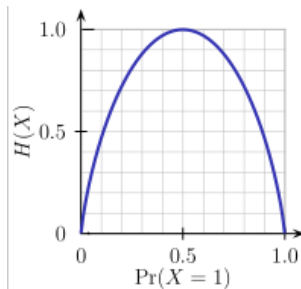
Here  $I$  is the information content of  $X$ .

## What is entropy(2)

When taken from a finite sample, the entropy can explicitly be written as

$$H(X) = \sum_i P(x_i) I(x_i) = - \sum_i P(x_i) \log_b P(x_i)$$

Common values of  $b$  are 2,  $e$ , and 10, and the unit of entropy is bit for  $b = 2$ , nat for  $b = e$ , and dit (or digit) for  $b = 10$ .



## Whis is entropy(3)

According the definition of  $H(X)$ , we can define **Joint Entropy** as:

$$H(X, Y) = - \sum_{x, y} p(x, y) \log p(x, y)$$

and define **Conditional Entropy** to measure unpredictability of random variable  $Y$  given random variable  $X$  as:

$$H(Y|X) = - \sum_{x, y} p(x, y) \log p(y|x)$$

The entropy or the amount of information revealed by evaluating  $(X, Y)$  (that is, evaluating  $X$  and  $Y$  simultaneously) is equal to the information revealed by conducting two consecutive experiments: first evaluating the value of  $Y$ , then revealing the value of  $X$  given that you know the value of  $Y$ . This may be written as

$$H(X, Y) = H(X|Y) + H(Y) = H(Y|X) + H(X) \quad (7)$$

## Whis is entropy(4)

We can get (7) by:

$$\begin{aligned} & H(X, Y) - H(X) \\ = & - \sum_{x,y} p(x, y) \log p(x, y) + \sum_x p(x) \log p(x) \\ = & - \sum_{x,y} p(x, y) \log p(x, y) + \sum_x \left( \sum_y p(x, y) \right) \log p(x) \\ = & - \sum_{x,y} p(x, y) \log p(x, y) + - \sum_{x,y} p(x, y) \log p(x) \\ = & - \sum_{x,y} p(x, y) \log \frac{p(x, y)}{p(x)} \\ = & - \sum_{x,y} p(x, y) \log p(y|x) = H(Y|X) \end{aligned}$$

## Whis is entropy(5)

Another useful measure of entropy that works equally well in the discrete and the continuous case is the **Relative Entropy** of a distribution. It is also called as the **Kullback-Leibler Divergence**.

$$D(p||q) = \sum_x p(x) \log \frac{p(x)}{q(x)} = E_{p(x)} \log \frac{p(x)}{q(x)}$$

To some extent, Relative Entropy can be the measure of the distance between two random variables.

**Mutual Information** can be defined as:

$$I(X, Y) = \sum_{x,y} \log \frac{p(x, y)}{p(x)p(y)}$$

# What is entropy(6)

Let's consider

$$\begin{aligned} & H(Y) - I(X, Y) \\ &= - \sum_y p(y) \log p(y) - \sum_{x,y} \log \frac{p(x, y)}{p(x)p(y)} \\ &= - \sum_y \left( \sum_x p(x, y) \right) \log p(y) - \sum_{x,y} p(x, y) \log \frac{p(x, y)}{p(x)p(y)} \\ &= - \sum_{x,y} p(x, y) \log p(y) - \sum_{x,y} \log \frac{p(x, y)}{p(x)p(y)} \\ &= - \sum_{x,y} p(x, y) \log \frac{p(x, y)}{p(x)} \\ &= - \sum_{x,y} p(x, y) \log p(y|x) = H(Y|X) \end{aligned}$$

According the formula (7):

$$I(X, Y) = H(X) + H(Y) - H(X, Y)$$